

Solution to Homework 2

Shoeb Mohammed and Zhuo Chen

February 7, 2016

1 Gradient and Hessian of $NLL(\theta)$ for logistic regression

1.1

Given

$$g(z) = \frac{1}{1 + e^{-z}} \quad (1)$$

Proof.

$$\frac{\partial g(z)}{\partial z} = \frac{(1 + e^{-z}) \cdot 0 + e^{-z}}{(1 + e^{-z})^2} = \frac{e^{-z}}{(1 + e^{-z})^2} = \frac{g(z)}{(1 - g(z))} \quad (2)$$

□

1.2

For logistic regression, negative log likelihood is

$$NLL(\theta) = - \sum_{i=1}^m (y_i \log(h_\theta(x_i)) + (1 - y_i) \log(1 - h_\theta(x_i))) \text{ where } h_\theta(x_i) = g(\theta^T x^i) \quad (3)$$

Proof. Using equation ?? and chain rule for differentiation we have

$$\begin{aligned} NLL(\theta) &= - \sum_{i=1}^m \left(\frac{y_i}{h_\theta(x_i)} h_\theta(x_i) (1 - h_\theta(x_i)) \frac{\partial \theta^T x^i}{\partial \theta} - \frac{(1 - y_i)}{(1 - h_\theta(x_i))} h_\theta(x_i) (1 - h_\theta(x_i)) \frac{\partial \theta^T x^i}{\partial \theta} \right) \\ &= - \sum_{i=1}^m (y_i (1 - h_\theta(x_i)) - (1 - y_i) h_\theta(x_i)) x_i \text{ because } \frac{\partial}{\partial \theta} \theta^T x^i = x^i \\ &= \sum_{i=1}^m (h_\theta(x_i) - y_i) \end{aligned} \quad (4)$$

□

1.3

Given

$$\begin{aligned} H &= X^T S X \text{ where } S = \text{diag}(\mu_1 \dots \mu_m) \\ \mu_i &= h_\theta(x_i)(1 - h_\theta(x_i)) \text{ for } i = 1 \dots m \\ \text{and } 0 < \mu_i < 1 &\text{ for } i = 1 \dots m \end{aligned} \quad (5)$$

Proof. For any vector $u \neq 0$ we have,

$$\begin{aligned} u^T H u &= u^T (X^T S X) u \\ &= (X u)^T S (X u) \\ &= v^T S v \text{ where } v = [v_1 \dots v_m]^T = X u \neq 0 \text{ since } X \text{ is full rank} \\ &= \sum_{i=1}^m v_i^2 \mu_i \\ &> 0 \text{ since } \mu_i \text{ is positive and } v_i \neq 0 \end{aligned} \quad (6)$$

Thus, H is positive definite. □

2 Regularizing logistic regression

Proof. The maximal likelihood and MAP estimates for θ are

$$\begin{aligned} \theta_{MLE} &= \underset{\theta}{\operatorname{argmax}} \prod_{i=1}^m P(y^{(i)} | x^{(i)}; \theta) \\ \theta_{MAP} &= \underset{\theta}{\operatorname{argmax}} P(\theta) \prod_{i=1}^m P(y^{(i)} | x^{(i)}; \theta) \text{ where } P(\theta) \sim N(0, \alpha^2 I) \end{aligned} \quad (7)$$

Equation ?? can be rewritten using log likelihood $LL(\theta)$:

$$\begin{aligned} \theta_{MLE} &= \underset{\theta}{\operatorname{argmax}} LL(\theta) \text{ where } LL(\theta) = \sum_{i=1}^m \log(P(y^{(i)} | x^{(i)}; \theta)) \\ \theta_{MAP} &= \underset{\theta}{\operatorname{argmax}} \log(P(\theta)) + LL(\theta) \\ &= \underset{\theta}{\operatorname{argmax}} K - \frac{d}{2\alpha^2} \theta^T \theta + LL(\theta) \text{ where } K \text{ is constant. This follows from } P(\theta) \sim N(0, \alpha^2 I) \\ &= \underset{\theta}{\operatorname{argmax}} LL(\theta) - \frac{d}{2\alpha^2} \|\theta\|_2^2 \end{aligned} \quad (8)$$

Now,

$$\begin{aligned}
LL(\theta_{MAP}) - \frac{d}{2\alpha^2} \|\theta_{MAP}\|_2^2 &\geq LL(\theta_{MLE}) - \frac{d}{2\alpha^2} \|\theta_{MLE}\|_2^2 && \text{from definition for } \theta_{MAP} \\
&\geq LL(\theta_{MAP}) - \frac{d}{2\alpha^2} \|\theta_{MLE}\|_2^2 && \text{from definition for } \theta_{MLE} \\
\implies \frac{d}{2\alpha^2} \|\theta_{MAP}\|_2^2 &\leq \frac{d}{2\alpha^2} \|\theta_{MLE}\|_2^2 \\
\implies \|\theta_{MAP}\|_2^2 &\leq \|\theta_{MLE}\|_2^2 \\
\implies \|\theta_{MAP}\|_2 &\leq \|\theta_{MLE}\|_2
\end{aligned} \tag{9}$$

□

3 Implementing logistic regression

3.1 Part A

Implementing logistic regression : the sigmoid function

Implemented sigmoid method in `utils.py`.

```
def sigmoid (z):
    sig = np.zeros(z.shape)
    # Your code here
    sig = 1/(1+np.exp(-1*z))
    # End your ode
    return sig
```

Cost function and gradient of logistic regression

Implemented loss and `grad_loss` methods in the `LogisticRegressor` class in `logistic_regressor.py`.

```
J=-1*np.sum(np.log( utils . sigmoid (np.dot(theta ,X.T)))*y+np.log(1- utils . sigmoid (np.dot(theta ,X.T)))*(1-y))/m
grad = np.dot(( utils . sigmoid (np.dot(theta ,X.T))-y) ,X)/m
```

Prediction using a logistic regression model

Implemented `predict` method in the `LogisticRegressor` class in `logistic_regressor.py`.

```
y_pred = utils . bin_features ( utils . sigmoid (np.dot(self.theta ,X.T))-0.5)
```

Add code in `ex1.py`.

```
pred_prob = utils . sigmoid (np.dot(log_reg1 . theta ,np.array ([1 ,45 ,85]).T))
accuracy = 1-float (np.sum(np.abs(predy-y))/y . shape [0])
```

Result


Run `ex1.py` in the shell:

```
Plotting data with green circle indicating (y=1) examples and red circle indicating (y=0) examples ...
Loss on all-zeros theta vector (should be around 0.693) = 0.69314718056
Gradient of loss wrt all-zeros theta vector (should be around [-0.1, -12.01, -11.26]) = [ -0.1          -12.00921659  -11.26284221]
Optimization terminated successfully.
Current function value: 0.203498
Iterations: 19
Function evaluations: 20
Gradient evaluations: 20
```


```

Theta found by fmin_bfgs: [-25.16056945  0.20622963  0.20146073]
Final loss = 0.203497702351
For a student with 45 on exam 1 and 85 on exam 2, the probability of admission = 0.776246678481
Accuracy on the training set = 0.89
Theta found by sklearn: [[-25.15293066  0.20616459  0.20140349]]

```



(a) The training data



(b) The decision boundary

Figure 1

All as expected.

3.2 Part B

Cost function and gradient for regularized logistic regression

Implemented loss and `grad_loss` methods in the `RegLogisticRegressor` class in `logistic_regressor.py`.

```

J=-1*np.sum(np.log(utils.sigmoid(np.dot(theta,X.T)))*y+np.log(1-utils.sigmoid(np.dot(theta,X.T))*(1-y))/m
+reg*np.sum(theta[1:]**2)/2/m

grad = np.dot((utils.sigmoid(np.dot(theta,X.T))-y),X)/m

```

Prediction using the model

Implemented `predict` method in the `LogisticRegressor` class in `logistic_regressor.py`.

```
y_pred = utils.bin_features(utils.sigmoid(np.dot(self.theta,X.T))-0.5)
```

Add code in `ex1_reg.py`.

```
accuracy = 1-float(np.sum(np.abs(predy-y))/y.shape[0])
```

Varying

Add code in `ex1_reg.py`.

```

for reg in [0,1.0,100.0]:
    ...

```

Exploring L1 and L2 penalized logistic regression

Add code in `ex1_reg.py`.

```
for reg in [0.1,0.5,1.0,5.0,10.0]:  
    ...
```

Result

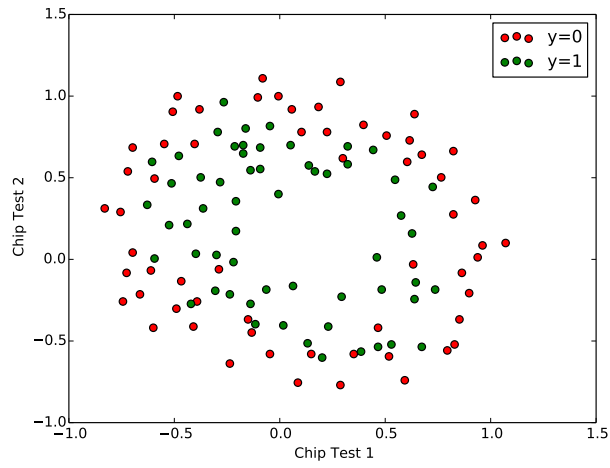
Run `ex1_reg.py` in the shell:

```
Plotting data with green circle indicating (y=1) examples and red circle indicating (y=0) examples ...  
Optimization terminated successfully.  
    Current function value: 0.224569  
    Iterations: 546  
    Function evaluations: 547  
    Gradient evaluations: 547  
Theta found by fmin_bfgs with reg = 0.0 [ 35.10191556  44.11916104  69.27187135 -344.27909285 -198.23463329  
-184.2284154 -295.82041756 -621.7326024 -510.84919909 -328.31173228  
1094.70040538 1269.58583356 1757.74907248  900.93789211  436.58879589  
471.12031352 1236.23835236 1822.81976807 1929.66695582 1131.05273288  
463.79908073 -1142.11739081 -2020.95888645 -3463.39935641 -3484.5099411  
-3252.26696072 -1546.00910831 -510.41253513]  
Final loss = 0.224568734073  
Accuracy on the training set = 0.915254237288  
Optimization terminated successfully.  
    Current function value: 0.529003  
    Iterations: 47  
    Function evaluations: 48  
    Gradient evaluations: 48  
Theta found by fmin_bfgs with reg = 1.0 [ 1.27268739  0.62557016  1.1809665 -2.01919822 -0.91761468 -1.43194199  
0.12375921 -0.36513086 -0.35703388 -0.17485805 -1.45843772 -0.05129676  
-0.61603963 -0.2746414 -1.19282569 -0.24270336 -0.20570022 -0.04499768  
-0.27782709 -0.29525851 -0.45613294 -1.04377851  0.02762813 -0.29265642  
0.01543393 -0.32759318 -0.14389199 -0.92460119]  
Final loss = 0.4624583499  
Accuracy on the training set = 0.830508474576  
Optimization terminated successfully.  
    Current function value: 0.621828  
    Iterations: 27  
    Function evaluations: 28  
    Gradient evaluations: 28  
Theta found by fmin_bfgs with reg = 5.0 [ 5.26750986e-01  8.29000343e-02  3.51747540e-01 -7.63457302e-01  
-2.16894856e-01 -4.73447564e-01 -6.09029000e-02 -1.03822986e-01  
-1.12858258e-01 -1.35090185e-01 -5.64095465e-01 -2.15435150e-02  
-2.05546999e-01 -5.63112422e-02 -4.64839657e-01 -1.56171750e-01  
-6.57923576e-02 -3.37041246e-02 -8.58443620e-02 -7.72930748e-02  
-2.70622817e-01 -4.14876764e-01 -1.61078926e-03 -1.01471865e-01  
2.46885978e-05 -1.10366999e-01 -2.49549621e-02 -4.24374938e-01]  
Final loss = 0.577292805563  
Accuracy on the training set = 0.813559322034  
Theta found by sklearn with L2 reg: with reg = 0.1 [ 2.65855183  1.76427994  2.91364412 -4.03385629 -3.34849756 -4.0181188  
0.76777199 -1.08648166 -0.47195071 -0.4774888 -3.27598952  0.54686285  
-1.80180787 -1.17932445 -2.79104067 -0.62127841 -0.4711418  0.61454641  
-1.14697992 -1.20796935 -0.10569617 -2.66246949  0.45857402 -0.76144039  
0.43744164 -1.17502213 -0.93753591 -1.20049576]  
Loss with sklearn theta: 0.353830932899  
Theta found by sklearn with L1 reg: [ 4.00212755  2.56718307  3.56329211 -7.68389893 -6.8113113 -8.66237057  
0.59188338 -0.20095689  0. 0. 2.445711 0.  
0. -1.7055534 0. 0. 0.36343109 -0.67166876 0.  
0. 0. -6.72105236 0. 0. 0. 0. 0.  
-0.06140466 0. ]  
Loss with sklearn theta: 0.33643326718  
Theta found by sklearn with L2 reg: with reg = 0.5 [ 1.57595698e+00  9.41010433e-01  1.64795490e+00 -2.54342145e+00  
-1.48161546e+00 -1.93049559e+00  2.76332771e-01 -5.72906659e-01  
-4.89795242e-01 -1.92532820e-01 -1.94870434e+00 -1.23689578e-02  
-8.94257808e-01 -4.62996758e-01 -1.60813528e+00 -3.02851338e-01  
-2.97201842e-01  1.63127043e-03 -4.40179153e-01 -4.71492765e-01  
-4.92042624e-01 -1.43161435e+00  8.76115308e-02 -4.19765000e-01  
5.79982479e-02 -4.93712970e-01 -2.66131793e-01 -1.15377891e+00]  
Loss with sklearn theta: 0.422346235529  
Theta found by sklearn with L1 reg: [ 2.74906778  1.55163984  2.19592256 -5.59295394 -3.43490014 -4.93205464  
0. 0. 0. 0. -2.27853613 0. 0.  
0. -2.11210668 0. 0. 0. 0. 0.  
0. 0. 0. 0. 0. 0. 0.  
0. ]  
Loss with sklearn theta: 0.37758866439  
Theta found by sklearn with L2 reg: with reg = 1.0 [ 1.1421394  0.60141117  1.16712554 -1.87160974 -0.91574144 -1.26966693  
0.12658629 -0.3686536 -0.34511887 -0.17368655 -1.42387465 -0.04870064  
-0.60646669 -0.26935562 -1.16303832 -0.24327026 -0.20702143 -0.04326335  
-0.28028058 -0.286921 -0.46908732 -1.03633961  0.02914775 -0.29263743  
0.01728096 -0.32898422 -0.13801971 -0.93196832]  
Loss with sklearn theta: 0.46843403006  
Theta found by sklearn with L1 reg: [ 1.86960916  0.68659701  1.2804614 -4.86238046 -1.62173321 -2.34246341  
0. 0. 0. 0. 0. 0. 0.  
0. -2.36718825 0. 0. 0. 0. 0.  
0. 0. 0. 0. 0. 0. 0.  
0. ]
```

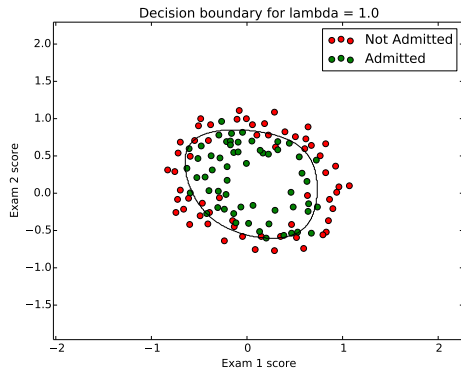
```

Loss with sklearn theta: 0.438149813414
Theta found by sklearn with L2 reg: with reg = 5.0 [ 4.01129749e-01  7.86245631e-02  3.59319093e-01 -6.87459023e-01
-2.25095774e-01 -3.94964659e-01 -5.66501493e-02 -1.02117116e-01
-1.07302805e-01 -1.26464125e-01 -5.35417213e-01 -2.40646947e-02
-1.94101075e-01 -5.88107997e-02 -4.39124052e-01 -1.52471333e-01
-6.51808095e-02 -3.18085564e-02 -8.46692847e-02 -7.44533648e-02
-2.68017771e-01 -4.01221135e-01 -2.79125965e-03 -9.76575762e-02
-6.04215794e-04 -1.07053968e-01 -2.56309837e-02 -4.16212768e-01]
Loss with sklearn theta: 0.58378743523
Theta found by sklearn with L1 reg: [ 0. 0. 0. -0.26092308 0. 0.
0. 0. 0. 0. 0. 0. 0.
0. 0. 0. 0. 0. 0. 0.
0. 0. 0. 0. 0. 0. ]
Loss with sklearn theta: 0.681052511162
Theta found by sklearn with L2 reg: with reg = 10.0 [ 0.21469236 -0.00761966 0.17611687 -0.4012903 -0.11745553 -0.23188083
-0.06668596 -0.05584267 -0.06215384 -0.09710193 -0.31766892 -0.01468057
-0.10913398 -0.03014551 -0.26764027 -0.11186999 -0.03627398 -0.02114738
-0.04753651 -0.04038118 -0.18117647 -0.24308692 -0.00364108 -0.05525352
-0.00101451 -0.06094026 -0.01293964 -0.26287463]
Loss with sklearn theta: 0.621592068026
Theta found by sklearn with L1 reg: [ 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0.
0. 0. 0. 0. 0. 0. 0. 0. 0. 0.]
Loss with sklearn theta: 0.69314718056


```



(a) The training data




(b) Training data with decision boundary of $\text{reg} = 1$ with $\text{sk}(12)$




(c) The regression path of $\text{reg} = 1$


Figure 2



(a) Overfitted boundary



(b) Properly fitted boundary



(c) Underfitted boundary

Figure 3: Boundary with different reg

All as expected. We can see overfitted and underfitted boundary when $\lambda=0$ and 100 respectively. When λ increases from 0.1 to 10 , the output of l_1 reg decreases to 0 more quickly than l_2 reg, and the cost of l_1 is higher than l_2 .

3.3 Part C

Feature transformation

Implemented `stdFeatures`, `logTransformFeatures` and `binarizeFeatures` method in `utils.py`.

```
def log_features(X):
    logf = np.zeros(X.shape)
    # Your code here
    logf = np.log(X+0.1)
    # End your code
    return logf
def bin_features(X):
    tX = np.zeros(X.shape)
    # your code here
    tX = np.array(X>0,dtype=int)
```



```
# end your code
return tX
```

Feature transformation

Implemented `select_lambda_crossval` method in `utils.py`.

```
def select_lambda_crossval(X,y,lambd_low,lambd_high,lambd_step,penalty):

    best_lambda = lambd_low

    # Your code here
    # Implement the algorithm above.
    best_accu=0
    kf=cross_validation.KFold(y.shape[0],10)
    l=lambd_low
    while l<=lambd_high:
        accu=0
        for train_index, test_index in kf:
            xx=X[train_index]
            yy=y[train_index]
            xt=X[test_index]
            yt=y[test_index]
            if penalty == "l2":
                lreg = linear_model.LogisticRegression(penalty=penalty,C=1.0/best_lambda, solver='lbfgs',fit_intercept=True)
            else:
                lreg = linear_model.LogisticRegression(penalty=penalty,C=1.0/best_lambda, solver='liblinear',fit_intercept=True)
            lreg.fit(xx,yy)
            predy = lreg.predict(xt)
            accu=accu+np.mean(predy==yt)
        accu=accu/10
        if accu>best_accu:
            best_accu=accu
            best_lambda=l
        l=l+lambd_step
    # end your code

    return best_lambda
```

Result

Run `ex1_spam.py` in the shell:

```
L2 Penalty experiments -----
best_lambda = 0.1
Coefficients = [-4.86311314] [[ -2.74144964e-02 -2.25297922e-01 1.21840741e-01 2.29363175e+00
2.70425757e-01 2.32851060e-01 9.28595406e-01 2.95200115e-01
1.62205894e-01 6.78255634e-02 -8.32602218e-02 -1.60373332e-01
-4.72248839e-02 1.07676572e-02 1.87904870e-01 8.19771733e-01
5.09529185e-01 3.98709124e-02 2.67729599e-01 3.47046562e-01
2.60498968e-01 3.64607069e-01 7.25020572e-01 1.96728174e-01
-3.15395736e+00 -4.03134022e-01 -1.25451015e+01 -6.16564148e-02
-1.56114501e+00 -5.51433441e-02 -3.00843039e-02 4.07264561e-01
-3.68156729e-01 -1.43612275e+00 -5.87186446e-01 4.44293905e-01
4.23160725e-02 -1.56897114e-01 -4.55330242e-01 -1.02250011e-01
-3.54273381e+00 -1.72944271e+00 -4.37530043e-01 -1.05999937e+00
-9.18599054e-01 -1.75490185e+00 -1.67475688e-01 -9.56877080e-01
-3.65654247e-01 -1.36535823e-01 -6.58693028e-02 2.06714195e-01
1.70694494e+00 1.21460221e+00 -3.35271575e-01 1.56142019e+00
3.68774998e-01]]
Accuracy on set aside test set for std = 0.9296875
best_lambda = 0.1
Coefficients = [-1.40695708] [[-0.1748124 -0.05135689 -0.12339923 1.2828822 0.46456529 0.26013964
0.96479518 0.48101469 0.11530886 0.15833164 -0.10051294 -0.1852999
-0.33761946 0.24045718 0.42765064 0.52228884 0.60140119 -0.10265903
0.12174456 0.29864954 0.20432531 0.18931018 0.54668794 0.56311458
-1.37213787 0.09174394 -3.52244763 0.25886851 -0.20749972 0.07686118
0.40469082 1.04149701 -0.15729015 1.29670296 -0.45104347 0.47004354
-0.35705768 0.36273978 -0.33026736 0.03533685 -0.31829402 -0.93876559
-0.68715245 -0.80497352 -0.44309598 -0.93099029 0.18540049 -0.84891725
-0.43476012 -0.12763685 0.19196783 0.7830564 1.47892381 0.02683296
0.69607044 0.04929847 0.362889482]]
Accuracy on set aside test set for logt = 0.943359375
best_lambda = 0.1
Coefficients = [-1.77183673] [[-0.27016277 -0.13652572 -0.44105224 0.19593277 1.09940287 0.28381492
2.39371265 0.89847243 0.26535465 0.435632 -0.39728324 -0.42882819
-1.07464275 0.32304905 0.66670341 1.58604221 1.10348806 -0.17358757
0.2491685 0.767151 0.78417213 1.4247451 1.03887036 1.61862715
-2.93738195 -0.23315908 -5.85002477 1.26265761 -0.83080782 -0.06708161
-1.65512023 -1.71956092 -0.92718438 0.44756264 -0.73938996 0.51818083
-1.06394957 1.07675048 -0.98242564 -0.32995306 -2.96965659 -2.5398783
-1.21092653 -2.16550138 -0.85215133 -2.56736788 0.03483552 -2.02741466
```

```

-0.37386944  0.22126778 -0.2152468  1.27339759  1.5956514  -0.03237894
-0.04463836 -0.04463836 -0.04463836]]
Accuracy on set aside test set for bin = 0.92578125
L1 Penalty experiments -----
best.lambda = 2.1
Coefficients = [-2.72840332] [[ -1.93809370e-02 -1.90165421e-01  1.26218367e-01  4.56664238e-01
 2.58423508e-01  2.02669188e-01  9.10326804e-01  2.92930859e-01
 1.58447578e-01  5.59232984e-02 -5.60502424e-02 -1.51431784e-01
-2.86214584e-02  1.08977734e-02  1.74617696e-01  7.93158066e-01
 4.83667987e-01  5.45908763e-02  2.63677888e-01  2.67168004e-01
 2.50977911e-01  3.57024814e-01  7.25613219e-01  2.17779610e-01
-2.75446443e+00 -3.78683120e-01 -6.94538791e+00 -3.94769408e-02
-6.47743850e-01 -9.90081163e-03  0.00000000e+00  0.00000000e+00
-3.58637375e-01  0.00000000e+00 -1.57522300e-01  3.32751667e-01
 4.09903100e-03 -1.34993641e-01 -3.79065669e-01 -6.46707542e-02
-6.35617521e-01 -1.10564315e+00 -2.72601780e-01 -7.83678731e-01
-8.24016676e-01 -1.45018868e+00 -1.23169859e-01 -7.13330756e-01
-3.01341655e-01 -1.31923358e-01 -5.95346647e-02  2.09510366e-01
 1.67669131e+00  5.59770230e-01 -7.64465943e-02  9.25879899e-01
 3.43024768e-01]]
Accuracy on set aside test set for std = 0.92578125
best.lambda = 5.1
Coefficients = [ 0.] [[-0.02237023  0. -0.05154852  0.16878365  0.42402059  0.15412008
 0.93326218  0.43544729  0.02545247  0.08676809  0. -0.18887613
-0.14290738  0.14542169  0.03694076  0.5082957  0.47578605  0.
 0.0795449  0.12335347  0.20266011  0.12228524  0.53917235  0.52854023
-1.08571406  0. -1.54138485  0.02033182  0. 0.
 0. -0.04969466  0. 0. 0.31653286 -0.36289535
 0. -0.11699187  0. 0. -0.55663169  0.
-0.35628821 -0.33859048 -0.77802727  0. -0.23452883 -0.02391159
 0. 0. 0.75476966  1.30466053  0. 0.55834284
 0.15493247  0.16648669]]
Accuracy on set aside test set for logt = 0.940755208333
best.lambda = 0.1
Coefficients = [-0.23362403] [[-0.27341794 -0.1210642 -0.43824938  0.13312448  1.10360255  0.29869675
 2.42028679  0.90527665  0.26731351  0.44562181 -0.41187711 -0.42707819
-1.09389899  0.3194352  0.66224276  1.59519397  1.13534114 -0.1819026
 0.24500927  0.80099193  0.78677071  1.41392778  1.03768696  1.65676144
-3.01060339 -0.19787025 -6.21759202  1.30722175 -0.83095457 -0.04058604
-1.74016234 -1.90105145 -0.9445934  0.20424823 -0.72605373  0.52532324
-1.0517627  1.10852783 -0.99529849 -0.30705931 -4.06519585 -2.61231263
-1.25451571 -2.21821251 -0.8609704 -2.58705802  0. -2.10652216
-0.3809371  0.22010838 -0.20768955  1.28165763  1.61412763 -0.0199922
-0.23270361 -1.0056229 -0.43930022]]
Accuracy on set aside test set for bin = 0.92578125

```

All as expected.

We can see that the l1 regularization could exclude unrelated features better. For 3 normalization methods, the logt method works the best.