

Solution to Homework 3

Shoeb Mohammed and Zhuo Chen

February 22, 2016

1 MAP and MLE parameter estimation

$\mathcal{D} = \{x^{(i)} | 1 \leq i \leq m\}$ where $x^{(i)} \sim \text{i.i.d } \text{Ber}(\theta)$

1.1

If m_1 are number of heads, m_0 are number of tails and $m_0 + m_1 = m$ then the likelihood and MLE for θ are

$$p(\mathcal{D}|\theta) = \theta^{m_1}(1 - \theta)^{m_0} \quad (1)$$

$$\begin{aligned} \theta_{MLE} &= \text{argmax}_{\theta} \theta^{m_1}(1 - \theta)^{m_0} \\ &= \text{argmax}_{\theta} m_1 \log \theta + m_0 \log(1 - \theta) \end{aligned} \quad (2)$$

θ_{MLE} satisfies (first derivative of the likelihood equals zero)

$$\frac{m_1}{\theta_{MLE}} - \frac{m_0}{1 - \theta_{MLE}} = 0 \quad (3)$$

Thus,

$$\begin{aligned} \theta_{MLE} &= \frac{m_1}{m_0 + m_1} \\ &= \frac{m_1}{m} \end{aligned} \quad (4)$$

1.2

The prior is

$$\begin{aligned} p(\theta) &= \text{Beta}(\theta|a, b) \\ &\propto \theta^{(a-1)}(1 - \theta)^{(b-1)} \end{aligned} \quad (5)$$

Thus, the posterior is

$$\begin{aligned} p(\theta|\mathcal{D}) &= \frac{p(\mathcal{D}|\theta)p(\theta)}{p(\mathcal{D})} \\ &= \frac{p(\mathcal{D}|\theta)p(\theta)}{\sum_{\theta'} p(\mathcal{D}|\theta')p(\theta')} \\ &\propto \theta^{m_1+a-1} \theta^{m_0+b-1} \end{aligned} \quad (6)$$

Thus,

$$\begin{aligned}\theta_{MAP} &= \operatorname{argmax}_{\theta} \theta^{m_1+a-1} \theta^{m_0+b-1} \\ &= \operatorname{argmax}_{\theta} (m_1 + a - 1) \log \theta + (m_0 + b - 1) \log(1 - \theta)\end{aligned}\tag{7}$$

Equation 7 is similar to MLE estimation. Thus,

$$\begin{aligned}\theta_{MAP} &= \frac{m_1 + a - 1}{m_0 + m_1 + a + b - 2} \\ &= \frac{m_1 + a - 1}{m + a + b - 2}\end{aligned}\tag{8}$$

It is clear from equations 8 and 4 that $\theta_{MAP} = \theta_{MLE}$ when $a = b = 1$.

2 Logistic regression and Gaussian Naive Bayes

2.1

$$\begin{aligned}p(y = 1|x) &= g(\theta^T x) \\ p(y = 0|x) &= 1 - g(\theta^T x)\end{aligned}\tag{9}$$

2.2

With naives Bayes assumption,

$$\begin{aligned}p(y = 1|x) &= \frac{p(x|y = 1)p(y = 1)}{p(x)} \\ &= \frac{p(x|y = 1)p(y = 1)}{p(x|y = 1)p(y = 1) + p(x|y = 0)p(y = 0)} \\ &= \frac{p(x|y = 1)\gamma}{p(x|y = 1)\gamma + p(x|y = 0)(1 - \gamma)} \quad \text{since } y \sim \operatorname{Ber}(\gamma) \\ &= \frac{\prod_{j=1}^d \mathcal{N}(\mu_j^1, \sigma_j^2)\gamma}{\prod_{j=1}^d \mathcal{N}(\mu_j^1, \sigma_j^2)\gamma + \prod_{j=1}^d \mathcal{N}(\mu_j^0, \sigma_j^2)(1 - \gamma)} \quad \text{since } p(x_j|y = 1) \sim \mathcal{N}(\mu_j^1, \sigma_j^2) \text{ and } p(x_j|y = 0) \sim \mathcal{N}(\mu_j^0, \sigma_j^2) \\ &= \frac{\mathcal{N}(\mu^1, \Sigma)\gamma}{\mathcal{N}(\mu^1, \Sigma)\gamma + \mathcal{N}(\mu^0, \Sigma)(1 - \gamma)} \quad \text{where } \mu_0 = (\mu_1^0 \cdots \mu_d^0)^T, \mu_1 = (\mu_1^1 \cdots \mu_d^1)^T, \Sigma = \operatorname{diag}(\sigma_1^2 \cdots \sigma_d^2)\end{aligned}\tag{10}$$

$$\begin{aligned}p(y = 0|x) &= \frac{p(x|y = 0)p(y = 0)}{p(x)} \\ &= \frac{p(x|y = 0)p(y = 0)}{p(x|y = 1)p(y = 1) + p(x|y = 0)p(y = 0)} \\ &= \frac{\mathcal{N}(\mu^0, \Sigma)\gamma}{\mathcal{N}(\mu^1, \Sigma)\gamma + \mathcal{N}(\mu^0, \Sigma)(1 - \gamma)} \quad \text{where } \mu_0 = (\mu_1^0 \cdots \mu_d^0)^T, \mu_1 = (\mu_1^1 \cdots \mu_d^1)^T, \Sigma = \operatorname{diag}(\sigma_1^2 \cdots \sigma_d^2)\end{aligned}\tag{11}$$

2.3

Proof. With uniform class priors, equation 10 gives

$$\begin{aligned}
p(y=1|x) &= \frac{\mathcal{N}(\mu^1, \Sigma)}{\mathcal{N}(\mu^1, \Sigma) + \mathcal{N}(\mu^0, \Sigma)} \\
&= \frac{1}{1 + \frac{\mathcal{N}(\mu^0, \Sigma)}{\mathcal{N}(\mu^1, \Sigma)}} \\
&= \frac{1}{1 + \frac{\exp(\frac{1}{2}(x-\mu^1)^T \Sigma^{-1}(x-\mu^1))}{\exp(\frac{1}{2}(x-\mu^0)^T \Sigma^{-1}(x-\mu^0))}} \\
&= \frac{1}{1 + \frac{\exp((x-\mu^1)^T \Lambda^2 (x-\mu^1))}{\exp((x-\mu^0)^T \Lambda^2 (x-\mu^0))}} \quad \text{where } \Lambda = \text{diag}\left(\frac{1}{\sqrt{2}\sigma_1} \cdots \frac{1}{\sqrt{2}\sigma_d}\right) \\
&= \frac{1}{1 + \frac{\exp((\Lambda(x-\mu^1))^T (\Lambda(x-\mu^1)))}{\exp((\Lambda(x-\mu^0))^T (\Lambda(x-\mu^0)))}} \\
&= \frac{1}{1 + \frac{\exp((\Lambda(z+a))^T (\Lambda(z+a)))}{\exp((\Lambda(z-a))^T (\Lambda(z-a)))}} \quad \text{where } a = \frac{\mu^0 - \mu^1}{2} \text{ and } z = x - \frac{\mu^0 + \mu^1}{2} \\
&= \frac{1}{1 + \exp((\Lambda(z+a))^T (\Lambda(z+a)) - (\Lambda(z-a))^T (\Lambda(z-a)))} \\
&= \frac{1}{1 + \exp(4(\Lambda a)^T (\Lambda z))} \\
&= \frac{1}{1 + \exp\left(4a^T \Sigma^{-1}\left(x - \frac{\mu^0 + \mu^1}{2}\right)\right)} \\
&= g(\theta^T x') \quad \text{where } \theta^T = [(\mu^0 - \mu^1)^T \Sigma^{-1}(\mu^0 + \mu^1), 2(\mu^1 - \mu^0)^T \Sigma^{-1}] \text{ and } x' = [1, x]
\end{aligned} \tag{12}$$

□

3 Softmax regression and OVA logistic regression

3.1 Implementing the loss function for softmax regression (naive version)

3.2 Implementing the gradient of loss function for softmax regression (naive version)

Implemented the `softmax_loss_naive` method in file `softmax.py`:

```

for i in range(0,m):
    p=np.zeros(max(y)+1)
    for j in range(0,max(y)+1):
        po=0
        for jj in range(0,max(y)+1):
            po=po+np.exp(theta[:,jj].dot(X[i,:])-theta[:,j].dot(X[i,:]))

```

```

p[j]=1/po
grad[:,j]=X[i,:]*(float(y[i]==j)-p[j])/m
J=J+np.log(p[y[i]])
J=-J/m+reg*np.sum(theta**2)/2/m
grad=grad+theta*reg/m

```

result:

```

Training data shape: (49000, 3072)
Validation data shape: (1000, 3072)
Test data shape: (10000, 3072)
Training data shape with bias term: (49000, 3073)
Validation data shape with bias term: (1000, 3073)
Test data shape with bias term: (10000, 3073)
loss: 2.33181510664 should be close to 2.30258509299
numerical: 1.846291 analytic: 1.846291, relative error: 1.620672e-08
numerical: 0.402461 analytic: 0.402461, relative error: 1.300510e-07
numerical: 2.983793 analytic: 2.983793, relative error: 9.064330e-09
numerical: 0.277037 analytic: 0.277037, relative error: 7.767378e-08
numerical: 1.066744 analytic: 1.066744, relative error: 5.981913e-08
numerical: -0.718366 analytic: -0.718366, relative error: 6.584340e-08
numerical: -0.298495 analytic: -0.298495, relative error: 1.193483e-07
numerical: 2.824531 analytic: 2.824531, relative error: 2.177955e-08
numerical: -0.617456 analytic: -0.617456, relative error: 1.193407e-08
numerical: 0.150777 analytic: 0.150777, relative error: 5.651458e-08

```

It performs as expected.

3.3 Implementing the loss function for softmax regression (vectorized version)

3.4 Implementing the gradient of loss function for softmax regression (vectorized version)

Implemented the `softmax_loss_vectorized` method in file `softmax.py`:

```

xt=X.dot(theta)
Pt=np.exp(xt-np.max(xt,1)).reshape([m,1]).dot(np.ones([1,theta.shape[1]]))
P=Pt/Pt.sum(1).reshape([m,1]).dot(np.ones([1,theta.shape[1]]))
J=-1.0/m*np.sum(np.multiply(np.log(P),convert_y_to_matrix(y)))+reg*np.sum(theta**2)/2/m
grad=-1.0/m*X.T.dot((convert_y_to_matrix(y)-P))+theta*reg/m

```

result:

```

naive loss: 2.331815e+00 computed in 2945.336793s
vectorized loss: 2.331815e+00 computed in 7.681536s

```

Loss difference: 0.000000
Gradient difference: 0.000000

we can see vectorized method is about 400 times faster then using for-loop because numpy has optimization for operating matrices, and it can get the same result.

3.5 Implementing mini-batch gradient descent

Implemented `train` and `predict` method of `SoftmaxClassifier` class in file `softmax.py`.

```
index=np.random.choice(range(0,len(y)),size=batch_size)
X_batch=X[index,:]
y_batch=y[index]

self.theta-=grad*learning_rate

y_pred=np.argmax(X.dot(self.theta),1)
```

3.6 Using a validation set to select regularization lambda and learning rate for gradient descent

3.7 Training a softmax classifier with the best hyperparameters

Codes for selecting best learning rate and regularization factor:

```
for lr in learning_rates:
for rs in regularization_strengths:
print("calculating: lr=ns=SoftmaxClassifier()
ns.train(X_train,y_train,lr,rs,verbose=True,batch_size=400,num_iters=2000)
ta=np.mean(y_train == ns.predict(X_train))
va=np.mean(y_val == ns.predict(X_val))
results[lr,rs]=(ta,va)
if va>best_val:
best_val=va
best_softmax=ns
```

We noticed that if λ is too large, the regularization term would be the major term in the gradient. That will cause the theta keep increasing with iteration. So we pick λ in the set $[5e4, 1e5, 5e5, 1e6, 5e6]$.

result:

```
lr 1.000000e-07 reg 5.000000e+04 train accuracy: 0.229163 val accuracy: 0.246000
lr 1.000000e-07 reg 1.000000e+05 train accuracy: 0.232490 val accuracy: 0.234000
lr 1.000000e-07 reg 5.000000e+05 train accuracy: 0.233694 val accuracy: 0.228000
lr 1.000000e-07 reg 1.000000e+06 train accuracy: 0.245714 val accuracy: 0.228000
```

```

lr 1.000000e-07 reg 5.000000e+06 train accuracy: 0.299347 val accuracy: 0.325000
lr 5.000000e-07 reg 5.000000e+04 train accuracy: 0.303429 val accuracy: 0.327000
lr 5.000000e-07 reg 1.000000e+05 train accuracy: 0.306898 val accuracy: 0.320000
lr 5.000000e-07 reg 5.000000e+05 train accuracy: 0.340245 val accuracy: 0.375000
lr 5.000000e-07 reg 1.000000e+06 train accuracy: 0.367102 val accuracy: 0.384000
lr 5.000000e-07 reg 5.000000e+06 train accuracy: 0.367959 val accuracy: 0.383000
lr 1.000000e-06 reg 5.000000e+04 train accuracy: 0.337286 val accuracy: 0.328000
lr 1.000000e-06 reg 1.000000e+05 train accuracy: 0.348980 val accuracy: 0.325000
lr 1.000000e-06 reg 5.000000e+05 train accuracy: 0.394367 val accuracy: 0.385000
lr 1.000000e-06 reg 1.000000e+06 train accuracy: 0.397102 val accuracy: 0.408000
lr 1.000000e-06 reg 5.000000e+06 train accuracy: 0.364531 val accuracy: 0.370000
lr 5.000000e-06 reg 5.000000e+04 train accuracy: 0.403082 val accuracy: 0.389000
lr 5.000000e-06 reg 1.000000e+05 train accuracy: 0.395612 val accuracy: 0.407000
lr 5.000000e-06 reg 5.000000e+05 train accuracy: 0.389490 val accuracy: 0.385000
lr 5.000000e-06 reg 1.000000e+06 train accuracy: 0.364673 val accuracy: 0.349000
lr 5.000000e-06 reg 5.000000e+06 train accuracy: 0.313286 val accuracy: 0.329000
best validation accuracy achieved during cross-validation: 0.408000
softmax on raw pixels final test set accuracy: 0.391400

```

An example of the iteration process(of the best classifier):

```

iteration 0 / 2000: loss 42.672658
iteration 100 / 2000: loss 24.842959
iteration 200 / 2000: loss 15.769520
iteration 300 / 2000: loss 10.091069
iteration 400 / 2000: loss 6.778380
iteration 500 / 2000: loss 4.788873
iteration 600 / 2000: loss 3.589836
iteration 700 / 2000: loss 2.930098
iteration 800 / 2000: loss 2.454714
iteration 900 / 2000: loss 2.182052
iteration 1000 / 2000: loss 2.055547
iteration 1100 / 2000: loss 1.979594
iteration 1200 / 2000: loss 1.913994
iteration 1300 / 2000: loss 1.864351
iteration 1400 / 2000: loss 1.880269
iteration 1600 / 2000: loss 1.748369
iteration 1700 / 2000: loss 1.846359
iteration 1800 / 2000: loss 1.843841
iteration 1900 / 2000: loss 1.864423

```

We noticed the loss starts to oscilate, which indicates 2000 iterations is enough.

the confusion matrix:

```

456 54 26 19 13 27 25 39 265 76
67 442 18 34 14 33 60 37 120 175
123 53 176 71 104 86 199 71 89 28
48 73 72 218 32 178 177 61 73 68

```

```

61 33 104 52 247 76 229 118 47 33
38 42 74 147 51 315 128 80 102 23
26 51 54 79 59 71 557 33 22 48
47 51 58 54 79 65 80 395 74 97
114 68 9 25 6 35 10 13 605 115
71 144 11 33 7 14 51 43 133 493

```

The (i,j) factor of this matrix means the count of image in class i that is predicted to be in class j , and numbers on the diagonal is the number of correctly predicted image. We can see that the performance of each class are different. the visualized theta is shown in figure 1.



Figure 1: the visualized theta